

✓

Congratulations! You passed!

TO PASS 80% or higher

Keep Learning

GRADE

100%

Week 3 Quiz

LATEST SUBMISSION GRADE

100%

1. If X is the standard notation for the input to an RNN, what are the standard notations for the outputs?

1 / 1 point

- ☐

Y
- ☐

H
- ☒

Y(hat) and H
- ☐

H(hat) and Y

✓

Correct

2. What is a sequence to vector if an RNN has 30 cells numbered 0 to 29

1 / 1 point

- ☐

The average Y(hat) for all 30 cells
- ☒

The Y(hat) for the last cell
- ☐

The Y(hat) for the first cell
- ☐

The total Y(hat) for all cells

✓

Correct

3. What does a Lambda layer in a neural network do?

1 / 1 point

- ☒

Allows you to execute arbitrary code while training
- ☐

Pauses training without a callback
- ☐

Changes the shape of the input or output data
- ☐

There are no Lambda layers in a neural network

✓

Correct

4. What does the axis parameter of tf.expand\_dims do?

1 / 1 point

- ☐

Defines the dimension index to remove when you expand the tensor
- ☒

Defines the dimension index at which you will expand the shape of the tensor
- ☐

Defines if the tensor is X or Y
- ☐

Defines the axis around which to expand the dimensions

✓

Correct

5. A new loss function was introduced in this module, named after a famous statistician. What is it called?

1 / 1 point

- ☐

Hyatt loss
- ☒

Huber loss
- ☐

Hubble loss
- ☐

Hawking loss

✓

Correct

6. What's the primary difference between a simple RNN and an LSTM

1 / 1 point

- ☒

In addition to the H output, LSTMs have a cell state that runs across all cells
- ☐

In addition to the H output, RNNs have a cell state that runs across all cells
- ☐

LSTMs have multiple outputs, RNNs have a single one
- ☐

LSTMs have a single output, RNNs have multiple

✓

Correct

7. If you want to clear out all temporary variables that tensorflow might have from previous sessions, what code do you run?

1 / 1 point

- ☐

tf.cache.backend.clear\_session()
- ☐

tf.keras.clear\_session
- ☐

tf.cache.clear\_session()
- ☒

tf.keras.backend.clear\_session()

✓

Correct

8. What happens if you define a neural network with these two layers?

1 / 1 point

tf.keras.layers.Bidirectional(tf.keras.layers.LSTM(32)),

tf.keras.layers.Bidirectional(tf.keras.layers.LSTM(32)),

tf.keras.layers.Dense(1),

- ☒

Your model will fail because you need return\_sequences=True after the first LSTM layer
- ☐

Your model will fail because you have the same number of cells in each LSTM
- ☐

Your model will fail because you need return\_sequences=True after each LSTM layer
- ☐

Your model will compile and run correctly

✓

Correct